



Experiment No. 7
Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:



**Aim:** Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model.

**Objective:** Able to perform various feature engineering tasks, perform dimensionality reduction on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

**Theory:**

In machine learning classification problems, there are often too many factors on the basis of which the final classification is done. These factors are basically variables called features. The higher the number of features, the harder it gets to visualize the training set and then work on it. Sometimes, most of these features are correlated, and hence redundant. This is where dimensionality reduction algorithms come into play. Dimensionality reduction is the process of reducing the number of random variables under consideration, by obtaining a set of principal variables. It can be divided into feature selection and feature extraction.

**Dataset:**

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.



## Vidyavardhini's College of Engineering & Technology

### Department of Computer Engineering

---

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

**Code:**



### **Conclusion:**

Comment on the impact of dimensionality reduction on the accuracy, precision, recall and F1 score.

#### **Accuracy:**

While dimensionality reduction can enhance accuracy in certain scenarios, excessive reduction may lead to information loss and a subsequent decrease in accuracy. In the case of the logistic regression model, the accuracy stands at approximately 0.8208.

#### **Precision:**

Dimensionality reduction has the potential to improve precision by reducing the occurrence of false positives. This reduction in noisy features allows the model to concentrate on the most valuable information. For the >50K class, both precision, recall, and the F1-score are 0.54, indicating an improvement in precision.

#### **Recall & F1 Score:**

In some instances, dimensionality reduction can boost recall by simplifying the data, making the model more effective at capturing underlying data patterns. For the class =50K, precision, recall, and the F1-score are 0.84, 0.94, and 0.89, respectively, highlighting an enhancement in recall and an overall balance in predictive performance.

```
import numpy as np
import pandas as pd
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
df=pd.read_csv('/content/adult.csv')
df.tail()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	0
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	0
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	0
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	0
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	0

```
df.describe().T
```

	count	mean	std	min	25%	50%	75%	max
age	32561.0	38.581647	13.640433	17.0	28.0	37.0	48.0	90.0
fnlwgt	32561.0	189778.366512	105549.977697	12285.0	117827.0	178356.0	237051.0	1484705.0
education.num	32561.0	10.080679	2.572720	1.0	9.0	10.0	12.0	16.0
capital.gain	32561.0	1077.648844	7385.292085	0.0	0.0	0.0	0.0	99999.0
capital.loss	32561.0	87.303830	402.960219	0.0	0.0	0.0	0.0	4356.0
hours.per.week	32561.0	40.437456	12.347429	1.0	40.0	40.0	45.0	99.0

```
df.shape
```

(32561, 15)

```
df.columns
```

Index(['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'capital.gain', 'capital.loss', 'hours.per.week', 'native.country', 'income'], dtype='object')

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 32561 entries, 0 to 32560  
Data columns (total 15 columns):  
# Column Non-Null Count Dtype  
--- -  
0 age 32561 non-null int64  
1 workclass 32561 non-null object  
2 fnlwgt 32561 non-null int64  
3 education 32561 non-null object  
4 education.num 32561 non-null int64  
5 marital.status 32561 non-null object  
6 occupation 32561 non-null object  
7 relationship 32561 non-null object  
8 race 32561 non-null object  
9 sex 32561 non-null object  
10 capital.gain 32561 non-null int64  
11 capital.loss 32561 non-null int64  
12 hours.per.week 32561 non-null int64  
13 native.country 32561 non-null object  
14 income 32561 non-null object  
dtypes: int64(6), object(9)  
memory usage: 3.7+ MB

```
df.isnull().sum()
```

age 0  
workclass 0  
fnlwgt 0  
education 0  
education.num 0  
marital.status 0  
occupation 0  
relationship 0  
race 0  
sex 0  
capital.gain 0  
capital.loss 0  
hours.per.week 0  
native.country 0  
income 0  
dtype: int64

```
for col in ['workclass','occupation','native.country']:  
    df[col].fillna(df[col].mode()[0],inplace=True)
```

```
df.isnull().sum()
```

age 0  
workclass 0  
fnlwgt 0  
education 0  
education.num 0  
marital.status 0  
occupation 0  
relationship 0  
race 0  
sex 0  
capital.gain 0  
capital.loss 0  
hours.per.week 0  
native.country 0  
income 0  
dtype: int64

```
df.replace({'Sex':{'male': 0, 'female':1}, 'Embarked':{'s': 0, 'C':1, 'Q':2}}, inplace=True)  
X = df.drop(['income'], axis=1)  
y = df['income']
```

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state = 0)
```

```
from sklearn import preprocessing
categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']
for feature in categorical:
    label = preprocessing. LabelEncoder()
    X_train[feature] =label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test [feature])
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = pd.DataFrame(scaler.fit_transform (X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
X_train.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	ca
0	0.101484	2.134215	-1.494279	-0.332263	1.133894	-0.402341	-0.600270	2.214196	0.39298	-1.430470	
1	0.028248	-1.279379	0.438778	0.184396	-0.423425	-0.402341	0.109933	-0.899410	0.39298	0.699071	
2	0.247956	0.086059	0.045292	1.217715	-0.034095	0.926666	-0.600270	-0.276689	0.39298	-1.430470	
3	-0.850587	-1.279379	0.793152	0.184396	-0.423425	0.926666	-0.363535	0.968753	0.39298	0.699071	
4	-0.044989	-1.962098	-0.853275	0.442726	1.523223	-0.402341	-0.600270	-0.899410	0.39298	0.699071	

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
LR= LogisticRegression()
LR.fit(X_train, y_train)
```

```
LogisticRegression
LogisticRegression()
```

```
y_pred=LR.predict(X_test)
accuracy_score(y_test, y_pred)
```

0.8203500870099294

```
from sklearn.decomposition import PCA
pca = PCA()
X_train = pca.fit_transform(X_train)
pca.explained_variance_ratio_

array([0.15112277, 0.10122703, 0.09056424, 0.0802928 , 0.07708238,
       0.07350038, 0.06774638, 0.06602885, 0.06115879, 0.06007244,
       0.05358847, 0.04835632, 0.04181168, 0.02744748])
```

```
X = df.drop(['income'], axis=1)
y=df['income']
```

```
X_train, X_test, y_train, y_test=train_test_split(X, y, test_size = 0.3, random_state = 0)
```

```
categorical=['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']
for feature in categorical:
    label = preprocessing. LabelEncoder()
    X_train[feature] =label.fit_transform(X_train[feature])
    X_test[feature]= label. transform(X_test [feature])
```

```
X_train= pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
```

```
pca= PCA()
pca.fit(X_train)
cumsum=np.cumsum(pca.explained_variance_ratio_)
dim=np.argmax(cumsum > 0.90) + 1
print('The number of dimensions required to preserve 90% of variance is',dim)
```

The number of dimensions required to preserve 90% of variance is 12

```
X = df.drop(['income', 'native.country', 'hours.per.week'], axis=1)
y = df['income']
X_train, X_test, y_train, y_test=train_test_split(X, y, test_size = 0.3, random_state = 0)
categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex']
for feature in categorical:
    label = preprocessing. LabelEncoder()
    X_train[feature]= label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test [feature])
```

```
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns=X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
```

```
LR2 = LogisticRegression()
LR2.fit(X_train, y_train)
```

```
LogisticRegression
LogisticRegression()
```

```
y_pred = LR2.predict(X_test)
accuracy_score (y_test, y_pred)
```

0.8208619101238612

```
from sklearn.metrics import confusion_matrix
import pandas as pd
confusion = confusion_matrix(y_test, y_pred)
df_confusion = pd.DataFrame (confusion, columns=['Predicted No', 'Predicted Yes'], index=['Actual No', 'Actual Yes'])
from sklearn.metrics import classification_report
print(classification_report (y_test, y_pred))
```

	precision	recall	f1-score	support
<=50K	0.84	0.94	0.89	7410
>50K	0.71	0.43	0.54	2359
accuracy			0.82	9769
macro avg	0.78	0.69	0.71	9769
weighted avg	0.81	0.82	0.80	9769